

DS 7347

# High-Performance Computing (HPC) and Data Science

## Session 8

---

Robert Kalescky

Adjunct Professor of Data Science

HPC Research Scientist

May 19, 2022

Research and Data Sciences Services

Office of Information Technology

Center for Research Computing

Southern Methodist University



Session Question

Docker

Singularity

Spack

Readings and Assignments

## Session Question

---



Why don't HPC systems use Docker?

# Docker

---



- There are several public and private sources for Docker images.
- Images can be used as the base image for custom images.
- Already optimized images can help with reproducible and efficient development workflows.



Docker <https://hub.docker.com>

Quay.io <https://quay.io>

NVIDIA <https://catalog.ngc.nvidia.com>

Intel <https://hub.docker.com/u/intel>

AMD <https://hub.docker.com/u/amdih>



```
1  #!/usr/bin/env sh
2
3  pull_and_check() {
4      name=${1}
5      tag=`echo ${name} | cut -d':' -f2`
6      docker pull ${name}
7      docker image ls | egrep "REPOSITORY|${tag}"
8      docker run --rm -it ${name} bash
9  }
10
11  images[0]="ubuntu:jammy"
12  images[1]="nvcr.io/nvidia/nvhpc:22.3-devel-cuda_multi-ubuntu20.04"
13  images[2]="nvcr.io/nvidia/nvhpc:20.7-runtime-cuda10.1-centos7"
14
15  for image in ${images[@]}; do
16      pull_and_check ${image}
17  done
18
```





- Images are CPU-architecture specific
- Docker supports multi-architecture builds
  - Platforms: amd64, arm32v5, arm32v6, arm32v7, arm64v8, i386, ppc64le, and s390x
  - `docker build --platform` with single platform
  - `docker buildx --platform` with list of platforms
- Builds on non-native platforms will be slower as it is running through a virtual machine



```
1 FROM ubuntu:20.04
2
3 ENV DEBIAN_FRONTEND noninteractive
4
5 RUN apt-get update &&\
6     apt-get -y install\
7     python3-pip\
8     python3-numpy\
9     python3-pandas
10
11 RUN pip3 install\
12     jupyterlab
13
14 ENTRYPOINT ["python3"]
15
```



```
1  #!/usr/bin/env sh
2
3  # Create builder to build images
4  docker buildx create --name builder --use
5
6  # Build images for x86_64 and ARM64
7  docker buildx build --platform\
8    linux/amd64,linux/arm64 -t rkalescky/python3:latest\
9    -f python3.dockerfile --push .
10
11 # Inspect the built images
12 docker buildx imagetools inspect rkalescky/python3:latest
13
```



- Images with build tools can be very large.
- Use the needed image for building.
- Use the smallest image for running.
- Define both the build and execution in a single Dockerfile.



```
1 FROM nvcr.io/nvidia/nvhpc:22.3-devel-cuda_multi-ubuntu20.04 as builder
2 WORKDIR /build
3 COPY hello_world.cpp ./
4 RUN nvc++ -Bstatic -o hello_world hello_world.cpp
5
6 FROM alpine:3.15.4
7 WORKDIR /opt/hello/bin
8 COPY --from=builder /build/hello_world ./
9 ENTRYPOINT ["/opt/hello/bin/hello_world"]
10
```



```
1  #!/usr/bin/env sh
2
3  # Build image using multi-stage Dockerfile
4  docker build -t hello:20.04 -f hello_world.dockerfile .
5
6  # Run the image
7  docker run hello:20.04
8
9  # Note the size difference
10 docker image ls | egrep "hello|22.3-devel"
11
```

# Singularity

---



- Singularity has it's own image **definition language**.
  - Requires (re)writing the definition file.
  - Requires root or “fakeroot”, which is not widely available on HPC systems.
  - Can be done on a Linux system with Singularity installed and then copying the image.
  - Not generally recommended as there would be two definition files to maintain, presumably Docker and also Singularity.
- Pull from Docker registries.
  - Requires pushing and pulling of Docker images.
- Build from Docker archives.
  - Requires exporting, copying, and conversion of Docker images.





```
1  #!/usr/bin/env sh
2
3  # Build images for both x86_64 and ARM64
4  . ./docker_buildx.sh
5
6  # Consume the image on M2 via Singularity
7  ssh m2 'bash -l -c "module load singularity\
8  && singularity run docker://rkalescky/python3 -c \'import numpy as np;
   ↪ print(np.pi)\'''
```



```
10 # Pull Docker image to a Singularity image
11 ssh m2 'bash -l -c "module load singularity\
12   && singularity pull -F python3_3.9.13-slim.sif docker://python:3.9.13-slim\
13   && ls -lh ./python3_3.9.13-slim.sif\
14   && singularity exec ./python3_3.9.13-slim.sif python3 -c \"import sys;
↪   print(sys.version)\""'
15
16 # Singularity mount points
17 ssh m2 'bash -l -c "module load singularity\
18   && echo $SINGULARITY_BIND"'
19
```



```
15  # Export, upload, convert, and run on M2 via Singularity
16  docker save python3:20.04 | ssh m2 'bash -l -c "n=python3_20_04\
17    && cat > ~/$n.tar\
18    && module load singularity\
19    && singularity build -F $n.sif docker-archive:$HOME/$n.tar\
20    && singularity exec $n.sif whoami"'
```

Spack

---



- Build images defined by Spack environments.
- Spack-based build optimizations are preserved.
- Intermediate Dockerfile uses multi-stage builds
- Currently does not work for multi-architecture builds.



```
1  spack:
2    specs:
3      - gromacs+mpi
4      - mpich
5    container:
6      format: docker
7      images:
8        os: "ubuntu:20.04"
9        spack: develop
10
```



```
1  #!/usr/bin/env sh
2
3  # Define the Spack environment
4  cat spack.yaml
5
6  # Build container definition file
7  spack containerize > Dockerfile
8
9  # Build the container image
10 docker build -t gromacs:latest .
11
```

## Readings and Assignments

---





## Readings

None

## Lab

- Use `spack containerize` to produce a Dockerfile of the Spack environment from Lab 1.
- Build and export the Docker image.
- Use Singularity on M2 to convert and run the image.
- Commit to your class repo:
  1. `assignments/lab_02.dockerfile`.
  2. `assignments/lab_02.{png,jpg}` of your terminal session showing a Python session with NumPy, Matplotlib, and Pandas loaded.
- Due 12:00 AM Central, Thursday, May 26, 2022