

DS 7347

# High-Performance Computing (HPC) and Data Science

## Session 10

---

Robert Kalescky

Adjunct Professor of Data Science

HPC Research Scientist

May 26, 2022

Research and Data Sciences Services

Office of Information Technology

Center for Research Computing

Southern Methodist University



Session Question

Slurm

- Resource Selection

- Dependencies

- Arrays

Readings and Assignments

## Session Question

---



Why do we use Slurm?

Slurm

---



There are many options for selecting specific resources.

- Nodes
- Cores
- Memory
- Process distribution
- Accelerators



```
1  #!/bin/bash
2  #SBATCH -J gpu
3  #SBATCH -o gpu_%j.out
4  #SBATCH -p development
5  #SBATCH --nodes=1
6  #SBATCH --ntasks-per-node=1
7  #SBATCH --cpus-per-task=1
8  #SBATCH --gres=gpu:1
9  #SBATCH --mem=6G
10
11  nvidia-smi
12
```



```
1  #!/bin/bash
2  #SBATCH -J mpi
3  #SBATCH -o mpi_%j.out
4  #SBATCH -p development
5  #SBATCH --nodes=3
6  #SBATCH --ntasks-per-node=1
7  #SBATCH --cpus-per-task=1
8  #SBATCH --mem=6G
9
10 # Build executable ahead of job submission
11 # module purge
12 # module load intel/oneAPI-2021
13 # mpicc -o mpi_get_hostnames mpi_get_hostnames.c
```





```
14
15  # Setup environment
16  module purge
17  module load intel/oneAPI-2021
18
19  # Run executable
20  srun mpi_get_hostnames
21
```



```
1  #!/bin/bash
2  #SBATCH -J mpi
3  #SBATCH -o mpi.out
4  #SBATCH -p development
5  #SBATCH --nodes=3
6  #SBATCH --ntasks-per-node=1
7  #SBATCH --cpus-per-task=2
8  #SBATCH --mem=6G
9  #SBATCH --exclude=k001
10
11 # Build executable ahead of job submission
12 # module purge
13 # module load intel/oneAPI-2021
14 # mpicc -fopenmp -o mpi_get_hostnames_hybrid mpi_get_hostnames_hybrid.c
```



```
16  # Setup environment
17  module purge
18  module load intel/oneAPI-2021
19  export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
20
21  # Run executable
22  srun mpi_get_hostnames_hybrid
23
```



- Jobs can be depended on others via the `--dependency` flag.
- There are numerous options for how this is done.
- `after` and `afterok` are common.



```
1  #!/usr/bin/env bash
2
3  for step in {1..10}; do
4      if [ ${step} -eq 1 ]; then
5          dependency_flag=""
6          previous_id="no job"
7      else
8          dependency_flag="--dependency=afterany:${dependency_id}"
9      fi
10     dependency_id=$(sbatch ${dependency_flag} --mem=1G --wrap "sleep 60")
11     dependency_id=${dependency_id[-1]}
12     printf "Job %s is dependent upon %s.\n" ${dependency_id} "${previous_id}"
13     previous_id=$(printf "job %s" ${dependency_id})
14 done
15
```



- Launch many similar jobs in parallel
- Execute over files, directories, data, etc. via shell scripting



```
1  #!/bin/bash
2  #SBATCH -J data
3  #SBATCH -o data_import_%A-%a.out
4  #SBATCH --array=0-2492           # Loops over all XML files
5  #SBATCH -p htc
6  #SBATCH --mem=6G
7  #SBATCH --cpus-per-task=1
8
9  # Load Python3
10 module purge
11 module load python
```



```
13  # Builds array of all original XML files, which are read-only
14  data_dir="$WORK/data/original_xml/"
15  readarray xml_files < <(ls ${data_dir}*.xml | sort)
16  f=${xml_files[${SLURM_ARRAY_TASK_ID}]}
17  fb=$(basename ${f} .xml)
18
19  # Converts XML to TSV and SQLite3 files
20  export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
21  python3 data_import.py ${f} ${fb}
22
```



## Readings and Assignments

---



## Readings

None



## Project

- Provide proposals for the following:
  1. Data source
  2. Analysis workflow
  3. Tools for implementing the workflow
  4. Possible performance optimization targets
- Commit to your class repo:
  1. `project/proposal.md`.
- Due 12:00 AM Central, Thursday, June 2, 2022